Purpose and Use Memo: Blockchain for Impact Funding Evidence Map

To: GSR Foundation

From: Michael Cooper, Research Director, Metagov (mike@metagov.org · Telegram: MCooperResearch)

> **Please note:** Everything in this memo is a draft and needs GSR's validation. It is based on GSR's initial answers about purpose, audience and success, and on GSR's public website. Proposed targets and design choices are suggestions, not commitments. The full Evidence Map depends on the outcome of a four-week pilot, described in Section 8. The final section lists the questions we need GSR to answer before the memo can be finalised.

1. Purpose

This memo records who the Blockchain for Impact Funding Evidence Map is for, what decisions it should inform, and how we will judge whether it was useful. Once agreed, it becomes the reference point for design choices and for each partner review milestone. If a proposed piece of work does not serve the uses below, we will drop it or raise it with GSR.

2. What the Evidence Map is

The Blockchain for Impact Funding Evidence Map is a structured, public overview of blockchain-based impact projects supported by GSR and its partners. For each project, it records what the project set out to achieve, how it was delivered, what happened after the pilot, and how strong the evidence is behind its reported results. It will show both where evidence exists and where it is missing.

The Evidence Map is designed to answer the following draft research questions:

1. What objectives are blockchain-based impact projects pursuing, and in which sectors and countries?

2. What implementation models do they use, including who delivers them, which blockchain features they rely on, and how they are funded?
3. How far do projects advance? Which ones never move beyond the pilot stage, and which are extended, scaled, handed over or shut down?
4. What is the quality of the evidence available on their performance? Does it cover only operational measures such as speed and cost, or does it include data on social outcomes?
5. How well did projects achieve their stated objectives, based on the available evidence?
6. Does any evidence show that blockchain added value compared with simpler, non-blockchain alternatives?
7. Where is evidence weak or missing, and which gaps matter most for funding decisions?

The unit of analysis is the use case, meaning one solution applied to one problem in one setting. Each use case will be classified by its objective, implementation model, blockchain function, deployment status and evidence grade, and by how well it achieved its objectives. The results will be presented as a matrix of use-case categories against outcomes, supported by a short report. The categories, definitions and grading rules will be finalized in the taxonomy developed during the pilot (Section 8).

3. Who the Evidence Map is for

The Evidence Map has two audiences, in order of priority.

1. **GSR and its partners.** This includes GSR staff and trustees, Mercy Corps Ventures, the UNICEF Venture Fund, the Filecoin Foundation, and GSR grantees that have received funding to implement blockchain-based solutions. Other partner organisations may be added later.
2. **The wider field.** This includes other funders, venture investors and implementing NGOs that are considering blockchain-based approaches.

All outputs will be fully public. Partners will review findings about their own use cases before publication, but no content will be withheld from the public version.

4. What the Evidence Map should inform

Primary use: targeting GSR's future Blockchain for Good funding. GSR needs to see where evidence is strong, where it is thin, and where new funding would close the most important gaps. GSR's Blockchain for Good stream accepts applications on a rolling basis rather than in fixed rounds (GSR How We Fund). The Evidence Map therefore needs to work as a standing reference for decisions, not as an input to a single funding round.

Secondary uses:

- **Positioning GSR as a convenor of evidence in the field.** GSR needs a credible, neutral base for bringing funders, investors and implementers together around shared evidence priorities.
- **Making the case to other donors.** GSR needs to show other funders, plainly and honestly, what blockchain-based approaches have and have not been shown to achieve, so they can decide whether to co-invest.

Supporting output: guidance on grantee reporting standards. The Evidence Map will identify what data future grantees would need to report to fill its gaps. The project will produce general guidance on what reporting standards could look like, so that GSR could later build them into its grantee design and submission requirements. Designing and putting those standards in place is not part of this project. However, the Evidence Map's categories and evidence grades will be built as a reusable data structure that future grantee reports could follow, so that new evidence adds to the map instead of starting again.

5. What the Evidence Map is not for

- **Evaluating individual grantees or partners.** The Evidence Map grades the evidence behind use cases. It does not judge the performance of any organisation.
- **Certifying projects.** The Evidence Map does not award or apply a kitemark. That would be the role of the possible follow-on project.

- **Proving that blockchain works.** The Evidence Map will report weak or negative evidence as clearly as positive evidence.
- **Replacing impact evaluations.** The Evidence Map shows where evaluations are needed. It does not produce new impact findings.
- **Setting binding grantee requirements.** The project offers guidance only. Any decision to adopt reporting standards rests with GSR.

6. How we will know it worked

The targets below are proposals for GSR to confirm or adjust. They apply only if the pilot in Section 8 leads to a decision to proceed with the full Evidence Map. Timeframes are measured from publication of the final report.

Use	Indicator	Proposed target	Timeframe
Targeting funding (primary)	GSR cites the Evidence Map in Blockchain for Good funding decisions or strategy documents	At least two funding decisions or one strategy document	Within 6 months
Grantee reporting guidance	GSR makes a decision on whether to develop reporting standards based on the guidance	Decision recorded	Within 6 months
Grantee reporting guidance	New grantee proposals or reports use the Evidence Map's categories and evidence grades	At least three proposals or reports	Within 12 months
Partner use	Partners use findings in their own strategy, portfolio or reporting decisions	At least two of the four partner groups	Within 12 months
UNICEF and Mercy Corps Ventures	Both organisations review the Evidence Map's categories and evidence grades and confirm whether they could use them in their own portfolio reporting	Response from both organisations	Within 6 months
Possible Kitemark	GSR decides whether to proceed with the Kitemark, using the Evidence Map's outputs as inputs	Decision recorded	Within 6 months
Convening	GSR presents the Evidence Map at funder or sector forums	At least two presentations	Within 6 months

Donor case-building	Another funder engages on the evidence agenda, for example through co-funding discussions or adopting the data structure	At least one funder	Within 12 months
Public good	Report, matrix, dataset and codebook are published openly	All four published	Within 4 weeks of completed partner review

7. What this means for the design

- **Gaps are framed as funding opportunities.** The gap analysis will be organised by use-case domain and outcome, so GSR can see directly where new grants would add the most evidence.
- **The categories are built to be reused.** The taxonomy, evidence grades and outcome fields will be documented as a standard that grantees, other funders and a possible Kitemark can apply.
- **Partners review before publication.** Because everything will be public and individual use cases will be named, partners need to know this from the start. They will check facts about their own use cases, while the research team keeps responsibility for evidence grades and conclusions.
- **The writing is for non-specialists.** The report will be written for funders and NGO leaders, with technical methods in annexes.
- **The Evidence Map can be updated.** The dataset and methods will be documented so the Evidence Map can be updated as new grantee data arrives.

8. Pilot Workplan

Before committing to the full Evidence Map, Metagov will run a four-week pilot using 20 to 30 public case studies from Mercy Corps Ventures, the UNICEF Venture Fund and the Filecoin Foundation. The pilot has two purposes: first, to test the methodology; and second, to

determine whether there is enough value in scaling up to the full exercise of 800+ use cases. GSR grantees will be added if the work moves to full scale. Michael Cooper will lead the analysis using an LLM. To test inter-rater reliability, meaning how consistently different coders grade the same material, two different LLMs will grade the same cases independently. The weeks below are relative, and the start date will be agreed with GSR.

Week 1: Pick the cases and agree the test

- Metagov will select 20 to 30 case studies from material that is already public: Mercy Corps Ventures/Prosper Global endline reports, UNICEF Venture Fund portfolio pages, and Filecoin Foundation case studies.
- The selection will aim for a spread across use cases, such as payments, identity, supply chain and data storage, and across outcomes, including some cases that stalled after the pilot.
- The pilot will use public sources only for now. This avoids the need for data-sharing agreements during the test.
- GSR and Metagov will agree the research question, including how to distinguish pilots from real-world deployment.

Week 2: Draft and apply the framework

- Metagov will draft a light taxonomy and the five-level evidence scale.
- Each case will also be given a deployment status: pilot, extended pilot, scaled or discontinued.
- Two different LLMs will grade the same cases independently.

Week 3: Test against the real world

- Metagov will compare the two LLMs' results and note where the taxonomy or scale broke down.

- For each case, Metagov will check what happened after the pilot: whether it is still running, scaled, handed over or shut down.
- Metagov will check whether the evidence grade predicts scaling. This is the "are pilots inflating utility?" question, and 20 to 30 cases are enough for a first read.
- Metagov will record which claims held up and which did not.

Week 4: Findings and decision

- Metagov will write a short memo covering what the cases show, where the method needs fixing, and whether the full 800+ exercise is worth doing.
- Metagov and GSR will agree whether to scale up, and what partner data access would add beyond the public material.
- The pilot findings memo will be shared only between GSR and Metagov.

Pilot metrics

The metrics below will help judge whether the method works and whether scaling up is worthwhile. The thresholds are proposals for GSR to confirm.

Metric	What it shows	Proposed threshold
1. Agreement between the two LLMs	Whether the evidence scale can be applied consistently, measured with a standard agreement statistic (Cohen's kappa)	0.7 or higher
2. Public-source coverage	Whether public material alone is enough to grade a case	80% or more of cases fully graded
3. Taxonomy fit	Whether the categories cover real cases	10% or fewer cases fit no category
4. Post-pilot status confirmed	Whether we can find out what happened after each pilot	75% or more of cases
5. Evidence grade and scaling	Whether cases with stronger evidence were more likely to scale	Reported as a first read, not pass or fail

6. Claims that held up	How far headline claims matched what happened after the pilot	Reported as a finding, not pass or fail

Proposed decision rule: We propose that the full Evidence Map should proceed if at least three of the four threshold metrics (1 to 4) are met and the method fixes identified in Week 3 can be made without redesigning the taxonomy or evidence scale. Metrics 5 and 6 show how much value a full exercise would add. GSR has final authority to approve or decline scaling up, based on its own criteria. (Amended 19 October 2026 with GSR's approval: the effort-per-case metric was removed.)

9. A possible follow-on: the Blockchain Funding Kitemark

GSR is interested in an effective kitemark, because blockchain project applications are difficult to understand. A possible follow-on project, the Blockchain Funding Kitemark, would answer one question: how can a trustee, international NGO or mainstream funder with no technical background judge whether a blockchain-based aid project is credible? It would produce a plain-language assessment framework covering governance design, key management (who controls the digital keys that move funds), regulatory compliance, dependence on "off-ramps" (services that convert digital currency into local cash), and whether a shared ledger is needed at all. It would also set a minimum bar for evidence. GSR's observation is that existing pilots report speed and cost, but almost none report social outcomes against what would have happened without the project. GSR would apply the framework to its own portfolio first, then offer it to UNICEF and Mercy Corps Ventures to co-sign. To stay light-touch, the framework would use plain-language questions a non-technical trustee can answer, take the form of a short checklist, and not require a technical audit. The Evidence Map would add value here without changing its primary purpose. Its categories, evidence grades and reporting guidance would be direct inputs, and its findings would test GSR's observation about current evidence so that the evidence bar is set at a realistic level. Whether the Evidence Map also records the Kitemark's assessment areas for each use case is still to be decided.

10. Team composition

Michael Cooper, Research Director at Metagov, will lead this effort for Metagov. He is an experienced social researcher and decision-support expert who has worked with the Millennium Challenge Corporation (MCC), USAID, the UK Foreign, Commonwealth and Development Office (FCDO) and the UNICEF Venture Fund. He led USAID's flagship platform for evidence gap maps and systematic reviews, giving him direct experience of the methods this project adapts. His work also includes leading an FCDO Frontier Tech Hub evaluation of blockchain pilots in agricultural insurance, land titling and humanitarian relief, as well as managing impact evaluations for MCC. This combination of evidence-synthesis leadership, evaluation of blockchain pilots and experience with the funders and partners involved makes him particularly well suited to this project.

Questions to be Answered

1. Do you agree with the audiences, primary use and secondary uses described in Sections 3 and 4?
2. Do the proposed targets and timeframes in Section 6 seem realistic? Which would you change?
3. What upcoming decisions would the Evidence Map inform, and when are they due? For example, trustee meetings, partnership reviews or strategy updates.
4. For each priority use, what finding from the Evidence Map would change what GSR does? For example, what would lead GSR to fund a domain it currently does not, or to stop funding one?
5. Are you comfortable with weak or negative evidence on GSR-funded use cases being published and named? Do partners need to agree to this in advance?
6. Besides Mercy Corps Ventures, UNICEF, the Filecoin Foundation and current grantees, which other stakeholders should be included as users or reviewers?
7. Who at GSR will be the main contact for the project, and who signs off on this memo?
8. Is anything missing from Section 5 that the Evidence Map should explicitly not be used for?

9. When should the pilot start, and who at GSR will take part in agreeing the research question in Week 1 and the scale-up decision in Week 4?
10. What criteria will GSR use to approve or decline scaling up to the full Evidence Map, beyond the pilot metrics in Section 8?
11. Do the proposed pilot metrics and thresholds in Section 8 seem right? Which would you change?
12. Are there specific public case studies GSR would like included in, or excluded from, the pilot?